

Project 2

Introduction to Computer Graphics

Andrzej Szymczak

October 25, 2013

Rendering the mesh

- ▶ Use the triangle soup mode (easiest to do)
- ▶ Remember you have to compute and specify normals for each vertex in order to get illumination working
- ▶ Normals are not provided with the input file
- ▶ Triangles are oriented consistently, so back-face culling should work
 - ▶ If back-face culling causes the back faces to be displayed by your implementation, just reorient your triangles (swap any pair of vertices, together with their attributes)

Mesh: Input file

- ▶ First two entries: triangle count T and vertex count V
- ▶ Triangle table (T triples of integers), sequence of triples of integer vertex IDs (in the range $0 \dots V - 1$) that define a triangle
- ▶ Vertex table coordinates of consecutive vertices
- ▶ Example: a tetrahedron

```
4 4
0 1 2
2 1 3
2 3 0
0 3 1
0.000000 0.000000 1.000000
1.000000 0.000000 0.000000
0.000000 1.000000 0.000000
0.000000 0.000000 0.000000
```

Vertex normals

- ▶ Vertex normals are typically computed as an average of normals of incident triangles
 - ▶ The process described below is tailored for smooth, high-resolution models (small triangles)
 - ▶ For smaller models, using the true geometric triangle normals may make more sense
- ▶ Normal of a vertex v = sum of cross products computed for all incident triangles
 - ▶ An area-weighted average
- ▶ How to compute the normals efficiently?

Vertex normals

- ▶ Linear-time algorithm:
 - ▶ Allocate a vector $N[]$ of V normals (V =number of vertices), initialize all entries to zero
 - ▶ For each triangle $\Delta(a, b, c)$ do
 - ▶ Below, a, b and c are considered integer IDs of vertices or their coordinates depending on context
 - ▶ Compute normal of Δ , $N_{\Delta} = \vec{ab} \times \vec{ac}$
 - ▶ $N[a] += N_{\Delta}$
 - ▶ $N[b] += N_{\Delta}$
 - ▶ $N[c] += N_{\Delta}$
- ▶ Since triangles are oriented consistently, all cross products will point in a consistent direction (i.e. will either be all outward or all inward)
- ▶ Want to use outward normals

Sending the mesh data to GPU

- ▶ In modern OpenGL, the information you need for rendering is stored on the GPU
- ▶ Need to upload arrays containing vertex (or other) data into a *buffer*
- ▶ Buffer is basically an array residing in the GPU memory
- ▶ GPU can access information in a buffer an order of magnitude faster than the same information in the main memory (e.g. 200GB/s vs 5GB/s)

Constructing the buffers

- ▶ First, build arrays containing vertex coordinates and normal vector coordinates in the main memory
 - ▶ It would also be possible to use a single array
 - ▶ Here we use one per attribute for simplicity
 - ▶ In most cases, you want buffers describing the same 3D model to contain the same number of attributes
 - ▶ Attributes could be of different sizes

Constructing the buffers

```
NormalArray = new GLfloat[3*verts]; // 3 floats per vert
CoordArray = new GLfloat[3*verts]; // for both attr's
i := 0;
for each triangle  $\Delta(a,b,c)$  do:
. NormalArray[3*i+0] = N[a].x;
. NormalArray[3*i+1] = N[a].y;
. NormalArray[3*i+2] = N[a].z;
. CoordArray[3*i+0] := va.x;
. CoordArray[3*i+1] := va.y;
. CoordArray[3*i+2] := va.z;
. ++i;
. do the same for vertex b...
. ++i;
. do the same for vertex c...
. ++i;
done
```


Sending buffers to GPU

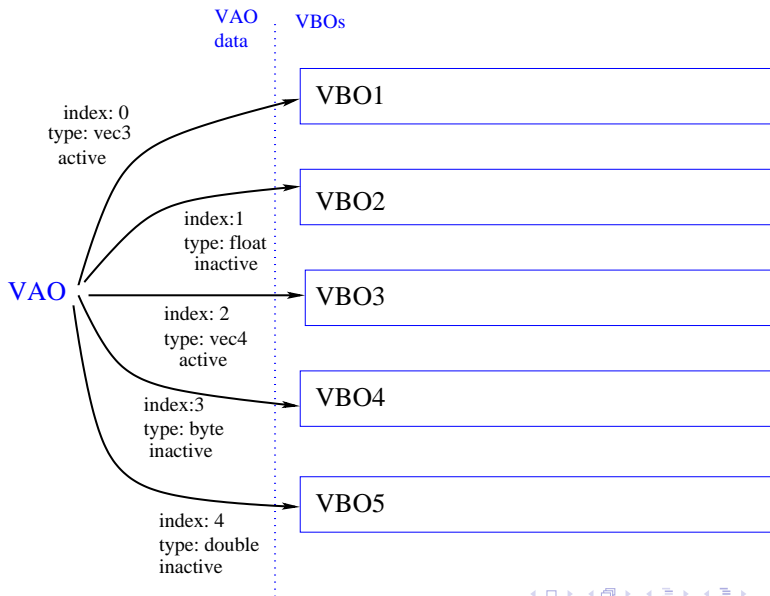
- ▶ Use the provided Buffer class for simplicity
- ▶ Remember:
 - ▶ A wrapper for an OpenGL Vertex Buffer Object (VBO)
 - ▶ No standard value/copy semantics implemented (usually no need for it)
 - ▶ Use pointers to Buffer objects in your code for best results
- ▶ Buffer Constructor
 - ▶ takes 3 arguments
 - ▶ number of *components* (basically, scalar values of the attribute) per entry; up to 4 allowed
 - ▶ number of entries
 - ▶ pointer to the array you want to put into the buffer; the code provides several overloaded versions to take care of different data types (note: no double/GLdouble!)
 - ▶ creates a buffer and sends array into it; stores the respective VBO information

```
Buffer *loc = new Buffer(3,verts,CoordArray);  
Buffer *normal = new Buffer(3,verts,NormalArray);
```

Vertex Array Objects (VAOs)

- ▶ A vertex typically has several different attributes
- ▶ Need a mechanism to tell the system which buffer to use to generate i -th attribute
- ▶ This is what VAOs are for
- ▶ `VertexArray`: a wrapper class for VAOs
- ▶ Construct an empty VAO using the default constructor
`myVAO = new VertexArray;`
- ▶ No copy/value semantics implemented - use pointers only
- ▶ Then, attach buffers to your VAO

VAOs and VBOs



Vertex Array Objects (VAOs)

- ▶ Attaching buffers to VAOs
- ▶ Use `attachAttribute` method

```
myVAO->attachAttribute(0,loc);
myVAO->attachAttribute(1,normal);
```
- ▶ If you use this VAO to generate vertices:
 - ▶ Attribute #0 will be looked up from the `loc` buffer (here: it will be location info)
 - ▶ Attribute #1 will be looked up from the `normal` buffer (for a given vertex, it will be the normal vector)
- ▶ Limitations of the OpenGL object wrappers
 - ▶ Don't use double precision attributes
 - ▶ Don't use matrix or array attributes
 - ▶ Don't mix floats with integer types (e.g. don't buffers with integer contents to specify values of floating point vertex attributes or the other way around)
 - ▶ `GLbyte`, `GLubyte`, `GLshort`, `GLushort`, `GLint`, `GLuint`, `GLfloat` attribute component types supported; in this project, `GLfloat` should be all you need

Vertex program

- ▶ C-like language (GLSL)
- ▶ Executed on the device (GPU)
- ▶ CPU–GPU communication needed for this to work!
- ▶ Input variables for the vertex shader: vertex attributes
 - ▶ Location qualifier provides the index information (this index has to match the index provided in the `attachAttribute` call)
 - ▶ For the VAO defined on previous slides, you would do something like this:

```
layout(location=0) in vec3 coord;  
layout(location=1) in vec3 normal;
```
- ▶ Now the system will know to put attribute #0 into `coord` and attribute #1 into `normal`

Vertex program: `gl_Position`

- ▶ Vertex shader assigns the projected vertex coordinates to the built-in output variable `gl_Position` of type `vec4`
- ▶ This provides information needed in the rasterization stage
 - ▶ Don't do division by the homogenous coordinate in the shader!
 - ▶ Typically, what you do is something like this:
`gl_Position = P * MV * coord;`
where `P` and `MV` are the projection and modelview matrices of type `mat4` and `coord` is a `vec4` containing vertex locations in the model coordinates
 - ▶ This can be done differently e.g. because:
 - ▶ World coordinates are needed for illumination calculations; you apply `MV` only, store the result in a variable, use it to evaluate the Phong's formula and then apply `P`
 - ▶ Some components of the modelview matrix may be represented differently (e.g. translation using a `vec3` vector, scaling using a `float` scale factor or rotation using a `mat3` instead using a 'full' `mat4` modelview matrix).

Converting vectors

- ▶ From standard to homogenous
 - ▶ Converting from `vec3` to homogenous coordinates:
`... vec3 coord;`
`vec4 coordh = vec4(coord,1.0);`
 - ▶ Similarly from `vec2` to homogenous coordinates (useful for flat shapes):
`... vec2 coord;`
`vec4 coordh = vec4(coord,0.0,1.0);`
- ▶ Dropping or reordering coordinates
 - ▶ Swizzle operations
 - ▶ Examples: `coordh.xyz`, `coordh.x`, `rgba.rgb`, `coordh.zyx`, `coordh.wwww`, etc
 - ▶ for one coordinate, the result is of type `float`; otherwise, it's `vecN`, with `N` equal to 2, 3 or 4
 - ▶ Can be used as l-values if no two repeated coordinate names

Vertex program: user defined output variables

- ▶ You can send some number of attributes with the processed vertex
- ▶ These attributes are interpolated on rasterization stage and send with fragments to fragment processing stage
- ▶ Declare an output variable using a line that contains:
 - ▶ interpolation qualifier (`flat`, `noperspective`, `smooth`);
`smooth` is the default and can be omitted;
`noperspective` = linear, `smooth` = perspective correct
 - ▶ key word `out`
 - ▶ type name (e.g. `float`, `int`, `vec3`, `vec4`...)
 - ▶ variable name; use the same name to declare the respective input to the fragment shader
- ▶ Examples of output variable declarations from the sample code
 - ▶ `out vec2 param;`
This one uses the default (perspective correct) interpolation
 - ▶ `flat out uint face_id;`
... and this one flat interpolation

Vertex program: user defined output variables

- ▶ Assign the values to all output variables in the shader's `main` function
- ▶ Other functions can be defined in GLSL code (probably no need to do that in out projects, but could be helpful for more complex shaders)
 - ▶ Functions may return values
 - ▶ Value-return convention: `in`, `out`, `inout` parameters

Fragment program

- ▶ Input variables need to match the output variables of the vertex shader
 - ▶ If the vertex shader has
`flat out uint face_id;`
then the fragment shader needs to have
`flat in uint face_id;`
 - ▶ If the vertex shader has
`out vec2 param;`
then the fragment shader needs to have
`in vec2 param;`
- ▶ In our projects, it is fine to declare a single output variable of type `vec3` and assign the final RGB values of the fragment to it in the `main()` function

Program objects

- ▶ The provided code makes dealing with them almost trivial
- ▶ Creation
 - ▶ Place the shader source code in two files (one for vertex, one for fragment shader)
 - ▶ Use the provided function `createProgram` to read, compile and link the shader code
 - ▶ Example from the sample code:

```
cube_program = createProgram("shaders/vsh_cube.glsl", "shaders/fsh_cube.glsl");
```

- ▶ Return value is of type `Program*`
 - ▶ **Don't use absolute paths!** or the grader will be mad...
- ▶ Turning on or off
 - ▶ `on` and `off` methods
 - ▶ `on` makes the program active, e.g. `cube_program->on()`;
 - ▶ `off` turns it off (more precisely: an invalid program whose result is undefined)
 - ▶ No memory; don't turn on two programs at once
- ▶ Deletion: use the provided destructor:
`delete cube_program;`

Program objects: uniform variables

- ▶ Uniform variables are technically not constant, but close to constant
 - ▶ they cannot be changed while a vertex stream is processed
 - ▶ ... but they can be altered (from the host code) between vertex streams
- ▶ Setting uniform variables
 - ▶ Use the `setUniform` method, for example
`cube_program->setUniform("T",-0.8f,0.8f);`
 - ▶ Arguments: name of the variable followed by values
 - ▶ Several overloaded versions available; examples:
`p->setUniform("A3Dvector",1f,1f,1f);`
`GLfloat v[4] = { 1.0, 2.0, 3.0, 1.0 };`
`p->setUniform("A4Dvector",v);`
 - ▶ The method should recognize how many components should be sent automatically, but it's best not to try to do anything that clearly does not make sense (e.g. send 4 numbers to a variable of type `vec2`)

Program objects: uniform variables


- ▶ General remarks on uniform variables
 - ▶ Each program object has a separate set
 - ▶ Shaders within a program object can share uniform variables; for example, if you declare
`uniform vec2 V;`
in both vertex and fragment shader, you can use it in both shaders (the value is the same in both and a single call to `setUniform` can be used to set it)

A few useful built-in GLSL functions

- ▶ `float dot (vec3 a, vec3 b)` computes the dot product of the arguments.
- ▶ `vec3 cross (vec3 a, vec3 b)` computes cross product.
- ▶ Scalar by vector multiplication, vector addition, subtraction and negation are available and work like overloaded operators. For example, if you have two variables `w` and `v` of type `vec3`, you can do `-w+0.8*v`.
- ▶ `vec3 normalize (vec3 a)` returns vector `a` normalized.
- ▶ Matrix by matrix multiplication and matrix by vector multiplication can be done by using overloaded `*` operators. Examples:

- ▶ For `M` of type `mat4`, i.e. 4×4 matrix, and a variable `v` of type `vec4`, you can do `M*v` to multiply them.
- ▶ There are types and operators for matrices of other dimensions (up to 4) available too – you may find `mat3` useful for storing normal transformation (equivalently, the rotational part of the modelview matrix).

Sample code output



Example: the cube program from sample code

```
#version 420

layout (location=0) in vec3 modelCoord;
layout (location=1) in uint faceId;

flat out uint face_id;

uniform mat4 MV;
uniform mat4 P;
uniform vec2 T;

void main()
{
    face_id = faceId;
    vec4 worldCoord = MV *
        vec4(0.4*(modelCoord-vec3(0.5,0.5,0.5))
            +vec3(T,0),1.0);
    gl_Position = P * worldCoord;
}
```

```
#version 420

flat in uint face_id;

out vec3 fragcolor;

void main()
{
    switch(face_id)
    {
        case 0: fragcolor = vec3(1,0,0); break;
        case 1: fragcolor = vec3(0,1,0); break;
        case 2: fragcolor = vec3(0,0,1); break;
        case 3: fragcolor = vec3(1,1,0); break;
        case 4: fragcolor = vec3(1,1,1); break;
        case 5: fragcolor = vec3(1,0,1); break;
        default: fragcolor = vec3(.5,.5,.5); break;
    }
}
```


Example: the square program from sample code

```
#version 420

layout (location=0) in vec2 model_coord;

out vec2 param;

uniform mat4 MV;
uniform mat4 P;

void main()
{
    vec3 scaled_coords = vec3(0.8*model_coord,0);
    param = 0.5*(model_coord+vec2(1,1));
    gl_Position = P * MV * vec4(scaled_coords,1);
}
```

```
#version 420

in vec2 param;

out vec3 fragcolor;

void main()
{
    float x = 512.0*param.x;
    float y = 512.0*param.y;

    fragcolor = vec3(
        0.5+0.5*sin(sin(x/30.0)+y*y/700.0),
        0.5+0.5*sin(y/71.0),
        0.5+0.5*sin(x*x*x/120000.0+y*y/1700.0)
    );
}
```

Sending vertices into pipeline

- ▶ Two ways to do this
 - ▶ 'plain': just look up values of attributes of consecutive vertices from the consecutive entries in the buffers attached to the VAO
 - ▶ indexed: use an index buffer to provide offsets for attribute lookup from the VAO
- ▶ Quite often, indexed saves memory (in particular, when some vertices need to be sent multiple times)
- ▶ Code for the plain method:

```
myVAO->sendToPipeline(GL_TRIANGLES,first,num);
```

 - ▶ First argument provides primitive setup control (GL_TRIANGLES means triangle soup; other values allowed: GL_TRIANGLE_FAN, GL_TRIANGLE_STRIP, GL_LINES, GL_LINE_STRIP, GL_POINTS)
 - ▶ 2nd argument: starting index
 - ▶ 3rd argument: how many vertices to form
 - ▶ attributes are looked up from the buffers attached to myVAO using offsets `first, first+1, ...first+num-1`

Sending vertices into pipeline: indexed

- ▶ Need an index buffer
- ▶ Use `IndexBuffer` class; as usual, use pointers only
- ▶ Construction: `myIx = new IndexBuffer(size, ptr);`
 - ▶ first argument: number of entries
 - ▶ second argument: pointer to an array containing the indices; must be of one of the unsigned types, `GLubyte`, `GLushort`, `GLuint`; of course, the array must be of size at least `size`
 - ▶ sends the data from the main memory to the GPU
- ▶ Code that creates the vertex stream:

```
myVAO->sendToPipelineIndexed(GL_TRIANGLES, myIx, first, num);
```

 - ▶ first argument: see previous slide
 - ▶ 2nd argument: pointer to `IndexBuffer` to be used
 - ▶ 3rd argument: index of the first vertex (into the index buffer)
 - ▶ 4th argument: number of vertices to be generated
 - ▶ Offsets used to generate vertex properties from buffers attached to the VAO: `myIx[first]`, `myIx[first+1]`, ... `myIx[first+num-1]`.

Sending vertices into the pipeline: example 1

```
GLfloat vertices[] = {
    -1, -1,
    -1,  1,
     1,  1,
     1, -1,
};
GLuint indices[] = {
    0, 1, 2,  // indices into the vertices of the first triangle
    0, 2, 3   // ... and second triangle (with consistent orientation)
};
.....
VertexArray *va_square = NULL;
Buffer *buf_square_vertices = NULL;
IndexBuffer *ix_square = NULL;
.....
buf_square_vertices = new Buffer(2,4,square_vertices);
ix_square = new IndexBuffer(6,square_indices);
va_square = new VertexArray;
va_square->attachAttribute(0,buf_square_vertices);
.....
va_square->sendToPipelineIndexed(GL_TRIANGLES,ix_square,0,6);
```

Example: the square programs

```
#version 420

layout (location=0) in vec2 model_coord;

out vec2 param;

uniform mat4 MV;
uniform mat4 P;

void main()
{
    vec3 scaled_coords = vec3(0.8*model_coord,0);
    param = 0.5*(model_coord+vec2(1,1));
    gl_Position = P * MV * vec4(scaled_coords,1);
}
```

```
#version 420

in vec2 param;

out vec3 fragcolor;

void main()
{
    float x = 512.0*param.x;
    float y = 512.0*param.y;

    fragcolor = vec3(
        0.5+0.5*sin(sin(x/30.0)+y*y/700.0),
        0.5+0.5*sin(y/71.0),
        0.5+0.5*sin(x*x*x/120000.0+y*y/1700.0)
    );
}
```

Sending vertices into the pipeline: example 2

```
VertexArray *va_cube = NULL;
Buffer *buf_cube_vertices = NULL;
Buffer *buf_cube_faceId = NULL;
```

```
.....
```

```
GLfloat cube_vertices[] = {
```

```
    0,1,0,  0,0,0,  0,0,1,
    0,1,0,  0,0,1,  0,1,1,
    0,0,1,  0,0,0,  1,0,0,
    0,0,1,  1,0,0,  1,0,1,
    0,0,0,  0,1,0,  1,1,0,
    0,0,0,  1,1,0,  1,0,0,
    1,0,1,  1,0,0,  1,1,0,
    1,1,0,  1,1,1,  1,0,1,
    1,1,0,  0,1,0,  0,1,1,
    0,1,1,  1,1,1,  1,1,0,
    1,1,1,  0,1,1,  0,0,1,
    0,0,1,  1,0,1,  1,1,1
```

```
};
```

```
GLubyte cube_faceId[] = {
```

```
    0,0,0,0,0,0,
    1,1,1,1,1,1,
    2,2,2,2,2,2,
    3,3,3,3,3,3,
    4,4,4,4,4,4,
    5,5,5,5,5,5
```

```
};
```

```
.....
```

```
buf_cube_vertices = new Buffer(3,36,cube_vertices);
buf_cube_faceId = new Buffer(1,36,cube_faceId);
va_cube = new VertexArray;
va_cube->attachAttribute(0,buf_cube_vertices);
va_cube->attachAttribute(1,buf_cube_faceId);
.....
va_cube->sendToPipeline(GL_TRIANGLES,0,36);
```

Example: the cube programs

```
#version 420

layout (location=0) in vec3 modelCoord;
layout (location=1) in uint faceId;

flat out uint face_id;

uniform mat4 MV;
uniform mat4 P;
uniform vec2 T;

void main()
{
    face_id = faceId;
    vec4 worldCoord = MV *
        vec4(0.4*(modelCoord-vec3(0.5,0.5,0.5))
            +vec3(T,0),1.0);
    gl_Position = P * worldCoord;
}
```

```
#version 420

flat in uint face_id;

out vec3 fragcolor;

void main()
{
    switch(face_id)
    {
        case 0: fragcolor = vec3(1,0,0); break;
        case 1: fragcolor = vec3(0,1,0); break;
        case 2: fragcolor = vec3(0,0,1); break;
        case 3: fragcolor = vec3(1,1,0); break;
        case 4: fragcolor = vec3(1,1,1); break;
        case 5: fragcolor = vec3(1,0,1); break;
        default: fragcolor = vec3(.5,.5,.5); break;
    }
}
```

- ▶ A header-only library that provides useful math functions
 - ▶ Vector operations
 - ▶ Matrix operations
 - ▶ Replicates functionality of most useful obsolete OpenGL functions such as transformation calls
 - ▶ Also, replicates most useful deprecated GL utilities (GLU) functions

Using GLM to set uniform matrix variables

- ▶ Setting the modelview and projection matrices
 - ▶ Suppose in the shader(s) of program prog you have:

```
uniform mat4 P;  
uniform mat4 MV;
```

- ▶ Example host code:

```
mat4 PMat = perspective(8.0f,1.0f,15.0f,25.0f);  
mat4 MVMat = translate(mat4(),vec3(0.0,0.0,-20.0)) *  
    rotate(mat4(),angle1,vec3(1.0,2.0,3.0)) *  
    rotate(mat4(),angle2,vec3(-2.0,-1.0,0.0)) *  
    scale(mat4(),vec3(0.1,0.1,0.1)) *  
    translate(mat4(),vec3(-c.x,-c.y,-c.z));  
prog->setUniform("P",&Pmat[0][0]);  
prog->setUniform("MV",&MVMat[0][0]);
```

Using GLM to set uniform matrix variables

- ▶ Setting the normal matrix

- ▶ In many cases (including this project) the normals are changed only by the rotational component – thus, we can use the 3×3 linear part of the modelview matrix as the normal matrix, or just use the rotation matrix only
- ▶ Using a plain rotation matrix could lead to slightly faster implementation of flat or Gouraud shading since rotation matrix is an isometry (and hence maps unit vectors into unit vectors)
- ▶ In general, you would take inverse and transpose of NMat used below before sending it to the uniform variable
- ▶ Host code:

```
mat3 NMat(MVMat);  
prog->setUniform("NM",&NMat[0][0]);
```


- ▶ Compatible vertex and/or fragment shader uniform declaration:

```
uniform mat3 NMat;
```

Setting up the transformations in project 2

- ▶ Problem: I am not telling you how big the model is and where it is "centered"
- ▶ Solution: scale and move it to a place that the camera can see!
- ▶ Modelview matrix and projection matrix are not entirely independent
- ▶ First step: "normalize" the model, i.e. scale and translate it so that it fits into a box extending from -1 to 1 in x , y and z

Model Normalization



- ▶ Use bounding box
- ▶ Compute $x_{min}, x_{max}, y_{min}, y_{max}, z_{min}, z_{max}$, where $?_{min}$ and $?_{max}$ are the minimum and the maximum $?$ -coordinates of the model's vertices for $? \in \{x, y, z\}$
- ▶ Translate the model by minus center of the bounding box
- ▶ Scale by $\frac{2}{\text{maximum dimension of the bounding box}}$

Model Normalization

- ▶ Translate the model by minus center of the bounding box,


$$\left(-\frac{x_{min} + x_{max}}{2}, -\frac{y_{min} + y_{max}}{2}, -\frac{z_{min} + z_{max}}{2} \right)$$

- ▶ Scale by $\frac{2}{\text{maximum dimension of the bounding box}}$, i.e. uniformly by

$$\frac{2}{\max(x_{max} - x_{min}, y_{max} - y_{min}, z_{max} - z_{min})}$$


- ▶ You may also want to apply rotation matrix (in the project, it will be specified using a trackball interface)
- ▶ Make sure order of transformations is correct
- ▶ Now it's time for the projection transformation...

Setting up projection



- ▶ First, pick your field of view (not too wide, something like 10-20 degrees max)
- ▶ In GLM angles (e.g. field of view in `glm::perspective` or rotation angle in `glm::rotate`) are specified in *degrees*; this means you'll likely have to do degree/radian conversion from time to time

Setting up projection



- ▶ First, pick your field of view α
- ▶ Your model (after normalization) is in the $-1...1$ box
- ▶ From trigonometry, $d = 1 / \tan(\alpha/2)$ (see figure)
- ▶ Translation by $(0, 0, -1 - d)$ places the $-1...1$ box in the field of view (good!)
- ▶ For distance to front/back clipping planes, one can use $d - 1$ and $d + 3$, to leave some slack for rotation (we don't want to clip the model even if it is rotated around its center)

Setting up transformations: summary

- ▶ Pick field of view α
- ▶ Modelview transformation:
 - ▶ Translate the model by minus center of the bounding box,


$$\left(-\frac{x_{min} + x_{max}}{2}, -\frac{y_{min} + y_{max}}{2}, -\frac{z_{min} + z_{max}}{2} \right)$$

- ▶ Scale by $\frac{2}{\text{maximum dimension of the bounding box}}$, i.e. uniformly by

$$\frac{2}{\max(x_{max} - x_{min}, y_{max} - y_{min}, z_{max} - z_{min})}$$

- ▶ You may also want to apply rotation matrix (in the project, it will be specified using a trackball interface)
 - ▶ Translate by $(0, 0, -1 - 1/\tan(\alpha/2))$
 - ▶ Arguments for `glm::project`: field of view α , aspect ratio 1 (square window), front and back clipping planes: $1/\tan(\alpha/2) - 1, 1/\tan(\alpha/2) + 3$

Project: trackball interface




- ▶ Imaginary sphere, projecting to circle inscribed into the window
- ▶ We'll pretend that parallel projection along z axis is used for the sphere and the sphere is centered at the origin and has unit radius

Project: trackball interface


- ▶ Say we are rendering into a $d \times d$ window
- ▶ Point on the sphere under a pixel (i, j) ?
- ▶ First, scale i and j to $[-1, 1]$ range (since the sphere is unit...)
 - ▶ $x := \frac{2i}{d-1} - 1$, $y := -(\frac{2j}{d-1} - 1)$ (since pixel coordinate y -axis points down)
- ▶ Now we have x and y coordinates of the point under the pixel (since sphere is projected parallel to z)
- ▶ If the viewpoint is at $(0, 0, +\infty)$, then $z = \sqrt{1 - x^2 - y^2}$
- ▶ The only complication is that the number under the root may be negative

Project: trackball interface



- ▶ This happens if the pixel is outside the projection of the sphere
- ▶ If this is the case, first project the pixel to the circle in a radial manner, then use the formula (beware of numerical issues though)


Project: trackball interface



- ▶ Radial projection: $B_0 = (x/\sqrt{x^2 + y^2}, y/\sqrt{x^2 + y^2})$
- ▶ Point under the pixel:

$$\begin{cases} (x, y, \sqrt{1 - x^2 - y^2}) & \text{if } 1 - x^2 - y^2 \geq 0 \\ (x/\sqrt{x^2 + y^2}, y/\sqrt{x^2 + y^2}, 0) & \text{otherwise} \end{cases}$$

Project: trackball interface



- ▶ Want rotation that takes p into q (p, q : depend on mouse movement)
- ▶ Axis: $\vec{cp} \times \vec{cq}$ where c is the center of the sphere
- ▶ Angle: $\angle(\vec{cp}, \vec{cq})$

Trackball interface: breakdown into event handlers

- ▶ Event-driven framework provided by the glut library
- ▶ Event handlers called not by you, but by glut (from within `glutMainLoop()` function, that basically dispatches events to handlers)
- ▶ Of course, you need to tell glut what they are `glutDisplayFunc`, `glutMouseFunc`, `glutMotionFunc`, `glutPassiveMotionFunc` etc.
- ▶ Useful global variables
 - ▶ R : the superposition of all 'finished' rotations (matrix)
 - ▶ R_0 : the rotation that is currently being specified using the trackball interface (matrix)
 - ▶ i_0, j_0 : coordinates of the last mouse button down event
- ▶ R and R_0 are initialized to identity at startup.
- ▶ Use glm's `mat3` or `mat4` type for R and R_0
- ▶ All the mouse event handlers get the coordinates of the mouse cursor (i, j) as arguments.

Trackball interface: breakdown into event handlers

- ▶ Mouse button down (begin dragging trackball): assign the event coordinates to i_0 and j_0 .
- ▶ Mouse moves with button down, current location: (i, j)
 - ▶ compute the points p and q on trackball under (i_0, j_0) and (i, j)
 - ▶ Set R_0 to the rotation that moves p into q ; be sure to use identity as that rotation if $i_0 = i$ and $j_0 = j$
 - ▶ Request redraw event by calling `glutPostRedisplay()`; want to give the user immediate feedback
- ▶ Mouse button up, current location: (i, j)
 - ▶ compute the points p and q on trackball under (i_0, j_0) and (i, j)
 - ▶ Let R' be the rotation that moves p into q or identity if $i_0 = i$ and $j_0 = j$
 - ▶ $R := R' * R$ (done dragging, so the rotation is finished)
 - ▶ $R_0 := Id$ (mouse button up; no rotation being specified)
- ▶ Drawing function
 - ▶ Apply rotation R and then rotation R_0 when drawing the model

Look at posts in the project directory

- ▶ Stray mouse button up events when selecting a menu entry
 - ▶ Select menu item by pressing the mouse button [generates menu event]
 - ▶ What goes down has to go up (at least a mouse button) [generates mouse button up event]
- ▶ Reading the input file
 - ▶ Keep it simple
 - ▶ Don't worry about the right number of white spaces etc; just read numbers